

Lab 1: Basic statistics in Python

Reference solution

Aymen Ben Brik
aymen.benbrik@esprit.tn
Esprit School of Business

2025–2026

Reference Python code is in `correction.py`; this document provides the analytical answers and expected numerical values.

Exercise 1 – Empirical moments

Solution

(1) For $n = 10\,000$ Gaussian draws (seed 42):

$$\hat{\mu} \approx 2.005, \quad \hat{\sigma}^2 \approx 3.987, \quad \hat{S} \approx -0.013, \quad \hat{K} \approx 2.996.$$

All four are close to their theoretical values $(2, 4, 0, 3)$, with a discrepancy of order $O(1/\sqrt{n})$ for the mean and variance.

(2) For $n = 100$, the variability of $\hat{\mu}$ is $\sigma/\sqrt{100} = 0.2$, so a typical deviation from $\mu = 2$ is of order 0.2. The plot of $\hat{\mu}$ vs n shows convergence toward μ at rate $1/\sqrt{n}$ (Law of Large Numbers).

Exercise 2 – Method of Moments vs MLE for the Gaussian

Solution

(1) For the Gaussian, $\hat{\mu}_{\text{MoM}} = \hat{\mu}_{\text{MLE}} = \bar{X}_n$ and $\hat{\sigma}_{\text{MoM}}^2 = \hat{\sigma}_{\text{MLE}}^2 = \frac{1}{n} \sum (X_i - \bar{X})^2$. The two functions return identical values on the same sample.

(2) On $B = 1000$ Monte-Carlo replicates with $n = 20$, $\sigma^2 = 4$:

$$\widehat{\text{Var}}_{\text{MLE}} \text{ mean} \approx 0.95 \sigma^2 = 3.80, \quad S_n^2 \text{ mean} \approx \sigma^2 = 4.00.$$

The factor $\frac{n-1}{n} = 0.95$ is exactly the bias predicted by theory.

Exercise 3 – Sampling distribution of $\bar{X}_n$

Solution

(1)–(2) The empirical variance of $\bar{X}_n$ over the B replicates is close to $\sigma^2/n = 4/n$:

n	σ^2/n	empirical variance
10	0.40	≈ 0.40
30	0.133	≈ 0.13
100	0.040	≈ 0.04
1000	0.004	≈ 0.004

The four histograms agree with the theoretical $\mathcal{N}(\mu, \sigma^2/n)$ density.

(3) For $\text{Exp}(1)$ ($\mu = \sigma^2 = 1$), the histogram of $\bar{X}_n$ is visibly skewed for $n = 10$ but well-approximated by $\mathcal{N}(1, 1/n)$ from $n \geq 30$ onwards — illustrating the CLT.

Exercise 4 – CI and t -test on Atlanta temperatures

Solution

The dataset `AvTempAtlanta.txt` contains 138 years (1879–2016) of monthly temperatures, stacked into $n = 138 \times 12 = 1656$ observations.

(1) Empirical statistics: $\bar{X}_n = 61.80$, $S_n = 13.37$. The 95% t -CI is

$$\text{CI}_{0.95}(\mu) = 61.80 \pm t_{1655, 0.975} \cdot \frac{13.37}{\sqrt{1656}} \approx 61.80 \pm 0.64 = [61.16, 62.45].$$

(2) For $H_0 : \mu = 60$ vs $H_1 : \mu \neq 60$:

$$T = \frac{61.80 - 60}{13.37/\sqrt{1656}} \approx 5.486, \quad t_{1655, 0.975} \approx 1.961.$$

$|T| \gg 1.961$, so we **reject** H_0 at the 5% level. The p -value is $\approx 4 \times 10^{-8}$. The mean monthly temperature is significantly different from 60°F.

(3) `scipy.stats.ttest_1samp(x, popmean=60)` returns the same statistic and p -value (up to floating-point precision).

Exercise 5 – Bonus: bootstrap CI

Solution

With $B = 5000$ resamples on the $n = 1656$ Atlanta observations:

$$\text{CI}_{0.95}^{\text{boot}}(\mu) \approx [61.15, 62.45],$$

essentially identical to the t -CI. Both methods agree because n is large (CLT very effective). The bootstrap would be more useful in a small-sample, non-Gaussian setting where the t -approximation is questionable.