

Lab 1: Basic statistics in Python

Moments · MoM · MLE · Inference

Aymen Ben Brik
aymen.benbrik@esprit.tn
Esprit School of Business

2025–2026

Goals

- Compute empirical moments and compare with theoretical values.
- Implement Method of Moments and MLE for the Gaussian.
- Study the sampling distribution of the sample mean by Monte-Carlo simulation.
- Build a confidence interval and run a t -test on a real dataset (Atlanta monthly temperatures).

Setup

- Python 3, dependencies: `numpy`, `scipy`, `matplotlib`.
- Reproducibility: `numpy.random.seed(42)`.
- Dataset: `data/AvTempAtlanta.txt` — monthly average temperatures, Atlanta, USA.
- The skeleton `seed.py` contains function signatures and the visualisation utilities.

Exercise 1 – Empirical moments

(1) Generate $n = 10\,000$ samples from $\mathcal{N}(2, 4)$ (mean 2, variance 4). Compute the four empirical moments:

$$\hat{\mu}, \quad \hat{\sigma}^2, \quad \hat{S} = \frac{1}{n\hat{\sigma}^3} \sum (X_i - \hat{\mu})^3, \quad \hat{K} = \frac{1}{n\hat{\sigma}^4} \sum (X_i - \hat{\mu})^4.$$

Compare with theoretical values $(\mu, \sigma^2, S, K) = (2, 4, 0, 3)$.

(2) Repeat for $n = 100$. Plot $\hat{\mu}$ as a function of n for $n \in \{10, 50, 100, 500, 1000, 5000, 10\,000\}$. What do you observe?

Exercise 2 – Method of Moments vs MLE for the Gaussian

(1) Implement `mom_gaussian(x)` and `mle_gaussian(x)` returning the estimated $(\hat{\mu}, \hat{\sigma}^2)$. Verify numerically that the two estimators *coincide* on the same sample.

(2) Compare the MLE/MoM estimator $\hat{\sigma}_{\text{MLE}}^2 = \frac{1}{n} \sum (X_i - \bar{X})^2$ with the unbiased one $S_n^2 = \frac{1}{n-1} \sum (X_i - \bar{X})^2$. Run $B = 1000$ Monte-Carlo replicates of size $n = 20$, draw the histograms of both estimators, and verify empirically that

$$\mathbb{E}[\hat{\sigma}_{\text{MLE}}^2] = \frac{n-1}{n} \sigma^2, \quad \mathbb{E}[S_n^2] = \sigma^2.$$

Exercise 3 – Sampling distribution of $\bar{X}_n$

- (1) For $n \in \{10, 30, 100, 1000\}$, simulate $B = 1000$ independent samples $X_1, \dots, X_n \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(\mu, \sigma^2) = \mathcal{N}(2, 4)$, and compute the sample mean of each.
- (2) Plot four histograms (one per n) of the B values of $\bar{X}_n$. Overlay the theoretical density of $\mathcal{N}(\mu, \sigma^2/n)$. Compute the empirical variance and compare with σ^2/n .
- (3) Repeat with a non-Gaussian distribution ($\text{Exp}(\lambda = 1)$, mean and variance both equal to 1). What does the CLT predict, and is it observed?

Exercise 4 – Confidence interval and t -test on Atlanta temperatures

Load `data/AvTempAtlanta.txt` (one column of monthly temperatures, n values).

- (1) Compute $\bar{X}_n$ and S_n . Build the 95% t -confidence interval for μ .
- (2) Test $H_0 : \mu = 60$ versus $H_1 : \mu \neq 60$ at level 5%. Report the test statistic, the critical value, the p -value, and the conclusion.
- (3) Compare with `scipy.stats.ttest_1samp`.

Exercise 5 – Bonus: bootstrap CI

- (1) Implement a non-parametric bootstrap CI for μ (resample the data $B = 5000$ times with replacement, compute $\bar{X}^*$ on each resample, take the empirical $(\alpha/2, 1 - \alpha/2)$ -quantiles of the bootstrap distribution).
- (2) Compare the bootstrap CI with the t -CI on the Atlanta data. Discuss when each is preferable.

Warning

Always set the seed (`np.random.seed(42)`) before any random draw. Save figures to PNG with explicit names (`lab1_ex3_sampling.png`, etc.).