

Project 2: Decomposition of monthly souvenir sales

Trend, seasonality and multiplicative model selection

Aymen Ben Brik
aymen.benbrik@esprit.tn
Esprit School of Business

2025–2026

Context

The dataset `data/SouvenirSales.csv` contains the monthly sales (in Australian dollars) of a souvenir shop on the beach of Queensland (Australia), from January 1995 to December 2001. The series exhibits both a strong upward trend and a marked seasonal pattern with a Christmas spike — a textbook example of a *multiplicative* time series.

Deliverables

- A reproducible Python script (`project2_<name>.py`, seed 42).
- A 4–6 page PDF report with figures, tables and discussion.
- Optional 5-minute oral defence.

Exercise 1 – Exploratory analysis

- (1) Load the data, parse the dates, plot Y_t vs. time.
- (2) Comment on:
 - the trend (is it roughly linear, exponential, or something else?);
 - the seasonal amplitude (constant or growing with the level?);
 - the presence of any obvious outliers.
- (3) Compute and report descriptive statistics (mean, std, min, max, skewness, kurtosis).

Exercise 2 – Additive vs. multiplicative

- (1) Apply an additive decomposition with `statsmodels.tsa.seasonal_decompose(y, period=12, model="additive")` and plot the four panels.
- (2) Same with `model="multiplicative"`.
- (3) Compare the residuals of the two models. Which is more homoscedastic?
- (4) Conclude: which model is more appropriate for this series, and why?

Exercise 3 – Log transform and additive workflow

- (1) Take $\log Y_t$. Plot the log-series.
- (2) On $\log Y_t$, apply the full additive workflow:
 - estimate the trend (LOESS, `frac = 0.3`, or polynomial degree 2);
 - detrend;

- estimate seasonality with seasonal averaging *and* a cosine-sine model with one harmonic;
- compute the residual.

(3) Diagnose the residual: histogram, QQ-plot, time-vs-residual plot. Is it stationary?

Exercise 4 – Forecasting

(1) Split the data: train = 1995-01 to 2000-12 (72 months), test = 2001-01 to 2001-12 (12 months).

(2) On the train set, fit a model $\log Y_t = \beta_0 + \beta_1 t + \sum_{j=1}^2 [a_j \cos(j\omega t) + b_j \sin(j\omega t)] + \varepsilon_t$ ($\omega = 2\pi/12$). Use OLS.

(3) Predict $\widehat{\log Y}_t$ for t in the test range. Convert back: $\widehat{Y}_t = \exp(\widehat{\log Y}_t)$.

(4) Compute the test MSE and MAPE. Plot observed vs. predicted on the test period.

Exercise 5 – Bonus: residual modelling

(1) Plot the autocorrelation function of the train residual (`statsmodels.graphics.tsaplots.plot_acf`).

(2) Are there significant autocorrelations beyond lag 0? This would suggest that the residual is not white noise — and motivates the ARMA modelling of Chapter 5.

Marking grid (out of 100)

Criterion	Points
Exploratory plots + descriptive statistics	15
Additive vs. multiplicative comparison + discussion	20
Log + full additive workflow + residual diagnostic	25
Forecast on test set: train/test split, OLS, MAPE	20
ACF of residual + ARMA motivation (bonus)	10
Reproducibility + report quality	10
Total	100

Aymen Ben Brik
aymen.benbrik@esprit.tn