

Lab 5: ARMA estimation, diagnostics and forecasting

Reference solution

Aymen Ben Brik
aymen.benbrik@esprit.tn
Esprit School of Business

2025–2026

Reference Python code in `correction.py`. All numerical values below are obtained with `numpy.random.seed(42)`.

Exercise 1 – AR(2) simulation and estimation

Solution

The PACF of the simulated series cuts off after lag 2, suggesting AR(2). Fit comparison:

order	log-lik.	AIC	BIC	LB(24) p
(1, 0, 0)	−720.28	1446.55	1459.19	0.0000
(2, 0, 0)	−684.91	1377.81	1394.67	0.7506
(3, 0, 0)	−684.84	1379.68	1400.76	0.7498

The MLE on (2, 0, 0) gives $\hat{\phi}_1 = +0.7051$, $\hat{\phi}_2 = -0.3645$, close to the true (0.6, −0.3) given $T = 500$. AR(1) is rejected by both AIC and Ljung–Box; AR(3) has a marginally larger BIC — AR(2) is the parsimonious choice.

Exercise 2 – MA(1) identification

Solution

Empirical $\hat{\rho}(h)$ for $h = 1, \dots, 5$:

$$(0.517, 0.024, -0.059, -0.073, -0.069),$$

the ACF cuts off at lag 1. The theoretical $\rho(1) = 0.7/1.49 = 0.4698$ is recovered to within ± 0.05 .

Fit comparison:

order	log-lik.	AIC	BIC	LB(24) p
(0, 0, 1)	−688.01	1382.02	1394.67	0.6726
(1, 0, 0)	−732.61	1471.22	1483.86	0.0000

$\hat{\theta} = +0.7558$ (true 0.7). The AR(1) alternative fails Ljung–Box because an MA(1) has only one non-zero autocorrelation, while AR(∞) is needed to reproduce that single spike via an autoregressive model: an AR(1) approximation leaves persistence on every lag, hence the residual autocorrelation.

Exercise 3 – ARMA(1, 1) on simulated data**Solution**

order	log-lik.	AIC	BIC	LB(24) p
(1, 0, 0)	−720.66	1447.31	1459.96	0.0000
(0, 0, 1)	−728.47	1462.93	1475.58	0.0000
(1, 0, 1)	−687.45	1382.91	1399.77	0.6568
(2, 0, 2)	−684.82	1381.64	1406.93	0.8687

AIC marginally favours (2, 0, 2), BIC favours (1, 0, 1). The (2, 0, 2) model adds three extra parameters for a 1.27-point gain in log-likelihood — classic over-fitting (AR/MA cancellation between the spurious factors). The parsimonious choice is (1, 0, 1) with $\hat{\phi} = +0.5093$ and $\hat{\theta} = +0.4730$, very close to the true (0.5, 0.4).

Exercise 4 – Box–Jenkins on the Atlanta residual**Solution**

On the Atlanta decomposition residual ($n = 1644$, $\text{sd} = 2.884$):

order	log-lik.	AIC	BIC	$\hat{\sigma}^2$	LB(24) p
(1, 0, 0)	−4071.62	8149.24	8165.46	8.293	0.0000
(2, 0, 0)	−4065.35	8138.70	8160.32	8.230	0.0000
(0, 0, 1)	−4071.22	8148.44	8164.65	8.289	0.0000
(1, 0, 1)	−3928.70	7865.39	7887.01	6.945	0.0000
(2, 0, 1)	−3871.89	7753.78	7780.81	6.480	0.0000

ARMA(2, 1) wins both AIC and BIC. The Ljung–Box test still rejects, however: the residual carries structure beyond what a linear ARMA can capture (residual seasonality / volatility clusters — to be addressed in Chapter 6).

Best parsimonious choice: **ARMA(2, 1)** as a useful approximation, with the caveat that the model is not a perfect fit.

Exercise 5 – Forecasting**Solution**

With ARMA(2, 1), the 24-step forecast (figure `lab5_ex5_forecast.png`) returns smoothly to the unconditional mean of zero, and the 95% confidence band widens toward

$$\pm 1.96 \hat{\sigma} \sqrt{\sum_{j=0}^{\infty} \psi_j^2} = \pm 1.96 \cdot (\text{unconditional sd of } X_t).$$

Hold-out validation. Splitting at 80% ($n_{\text{train}} = 1315$, $n_{\text{test}} = 329$): out-of-sample RMSE = 2.699, to be compared with the in-sample residual $\text{sd} = 2.598$. The out-of-sample error is only $\approx 4\%$ larger than the in-sample sd : the model generalises reasonably well, but the seasonal / volatility structure noted in Exercise 4 shows up as small systematic forecast errors.