

Final integrative project  
End-to-end time-series pipeline  
Statistics · Decomposition · ARMA · GARCH · VaR

Aymen Ben Brik  
aymen.benbrik@esprit.tn  
Esprit School of Business

2025–2026

## Context and objectives

This integrative project mobilises the notions of every chapter of the module:

- **Ch. 1 Basic statistics:** descriptive analysis, confidence intervals, normality tests.
- **Ch. 2 Decomposition:** classical time-series decomposition.
- **Ch. 3 Stationarity:** ACF/PACF and Ljung–Box.
- **Ch. 4 Non-stationarity:** ADF/KPSS unit-root tests and differencing.
- **Ch. 5 ARMA:** Box–Jenkins identification, MLE, forecasting.
- **Ch. 6 ARCH/GARCH:** Engle LM, GARCH(1, 1), and Value-at-Risk.

The fil-rouge task is a complete *risk-management study* of a synthetic but realistic daily stock-price series. The goal is to produce, by the end of the project, a one-day Value-at-Risk that is well-calibrated according to a Kupiec back-test.

## Setup

- Language: Python 3, dependencies: `numpy`, `pandas`, `matplotlib`, `scipy`, `statsmodels`, `arch`.
- Reproducibility: `numpy.random.seed(42)`.
- Skeleton in `seed/projet.py`; full reference solution in `correction/projet.py`.
- Dataset: `data/stock_prices.csv` –  $T = 2000$  synthetic daily prices with drift and GARCH volatility, regenerated by `generate_dataset.py`.
- Deliverables:
  - a reproducible `projet_<name>.py`;
  - a 10–12 page report PDF with figures and discussion;
  - a 12-min individual oral defence.

## Data

The data file `data/stock_prices.csv` contains  $T = 2000$  daily closing prices  $P_t$  over the period 2018-01-02 to 2025-09-15 (business days). The series is generated as

$$\log P_t = \log P_{t-1} + \mu + \varepsilon_t, \quad \varepsilon_t = \sigma_t z_t, \quad z_t \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, 1),$$

with  $\sigma_t^2 = \alpha_0 + \alpha_1 \varepsilon_{t-1}^2 + \beta_1 \sigma_{t-1}^2$  – a textbook GARCH(1, 1) data-generating process on the log-returns. The true parameters are kept hidden; you will recover them from the data.

### Exercise 1 – Part A – Basic statistics on returns (Ch. 1)

- (A.1) Load `data/stock_prices.csv`. Compute the log-returns  $r_t = \log(P_t/P_{t-1})$ ,  $t = 1, \dots, T-1$ .
- (A.2) Report the sample mean  $\hat{\mu}$ , sample standard deviation  $\hat{\sigma}$ , sample skewness, sample excess kurtosis. Provide a 95% confidence interval for the mean (using the Student- $t$  approximation).
- (A.3) Test the null hypothesis  $H_0 : \mu = 0$  at the 5% level. Conclude.
- (A.4) Plot a histogram of  $r_t$  together with the  $\mathcal{N}(\hat{\mu}, \hat{\sigma}^2)$  density. Run a Jarque–Bera normality test and comment.

### Exercise 2 – Part B – Decomposition of the price level (Ch. 2)

- (B.1) On the price series  $P_t$  (*not* the returns), extract the trend by a centred moving average of width  $k = 50$  business days.
- (B.2) On the log-price  $\log P_t$ , fit a linear trend by OLS:

$$\log P_t = \beta_0 + \beta_1 t + u_t.$$

Report  $\hat{\beta}_1$  – the implied mean log-return. Compare with the sample mean of  $r_t$  (Part A).

- (B.3) Plot  $P_t$ , the moving-average trend and the OLS exponential trend on the same figure. Comment in 5 lines.

### Exercise 3 – Part C – Stationarity diagnostics (Ch. 3)

- (C.1) Plot the sample ACF of  $\log P_t$  up to lag 60. What do you observe? Is the series consistent with a stationary process?
- (C.2) Same on the returns  $r_t$ . Are returns serially uncorrelated?
- (C.3) Run the Ljung–Box test on  $r_t$  at lags  $m \in \{10, 20\}$ . Conclude.
- (C.4) Run the Ljung–Box test on  $r_t^2$  at the same lags. What do the two results, taken together, suggest?

### Exercise 4 – Part D – Unit-root testing (Ch. 4)

- (D.1) Run the ADF test on  $\log P_t$  with `regression="ct"` (constant + trend) and on  $r_t$  with `regression="c"`. Report stat,  $p$ -value, decision at 5% in a small table.
- (D.2) Run the KPSS test on the same two series.
- (D.3) Build the ADF/KPSS combination matrix on  $\log P_t$ . In which cell does it land? What does this imply about the integration order  $I(d)$  of the price series?
- (D.4) Confirm:  $\log P_t \sim I(1)$  and  $r_t \sim I(0)$ .

### Exercise 5 – Part E – ARMA modelling (Ch. 5)

- (E.1) On the returns  $r_t$ , plot the sample ACF and PACF up to lag 20. Identify any structure in the conditional mean.
- (E.2) Fit three candidates by MLE (`statsmodels.tsa.arima.ARIMA`):

- ARMA(0,0): constant mean only;
- ARMA(1,0);
- ARMA(0,1).

Report a comparison table: log-likelihood, AIC, BIC, Ljung–Box(24)  $p$ -value.

**(E.3)** Choose the best model on the (AIC, BIC, Ljung–Box) triplet. The constant-mean model often wins – justify why this is expected for daily returns.

### Exercise 6 – Part F – GARCH(1,1) and VaR (Ch. 6)

**(F.1)** On the residuals  $\hat{\varepsilon}_t$  of the chosen mean model from Part E, run Engle’s LM test at lags 5, 10, 20. Confirm the presence of ARCH effects.

**(F.2)** Fit a GARCH(1,1) on the returns  $r_t$  with a constant mean and Gaussian innovations using the `arch` package. Report  $\hat{\mu}$ ,  $\hat{\alpha}_0$ ,  $\hat{\alpha}_1$ ,  $\hat{\beta}_1$  with their standard errors. Compute  $\hat{\alpha}_1 + \hat{\beta}_1$  and the implied unconditional variance.

**(F.3)** Compute the standardised residuals  $\hat{z}_t = \hat{\varepsilon}_t / \hat{\sigma}_t$ . Run Engle’s LM test on  $\hat{z}_t^2$ . Has the GARCH absorbed the volatility clustering?

**(F.4)** Compute the conditional one-day Value-at-Risk

$$\text{VaR}_{T+1}(\alpha) = -(\hat{\mu} + z_\alpha \hat{\sigma}_{T+1}), \quad \alpha \in \{1\%, 5\%\}.$$

Compare with the unconditional VaR (using the sample  $\hat{\sigma}$ ).

### Exercise 7 – Part G – Kupiec back-test (synthesis)

**(G.1)** Split the data: train = first 1750 obs, test = last 250 obs.

**(G.2)** On the test set, compute a *rolling* conditional VaR at  $\alpha = 5\%$ : at each day  $t$  in the test set, fit the GARCH on observations  $1..t-1$ , compute  $\hat{\sigma}_t$ , and predict  $\text{VaR}_t(5\%)$ .

**(G.3)** Count the number of *exceedances* – days where the realised loss  $-r_t$  exceeded  $\text{VaR}_t(5\%)$ .

**(G.4)** Run the Kupiec proportion-of-failures test:

$$LR_{\text{Kupiec}} = -2 \log \left[ \frac{(1-\alpha)^{N-X} \alpha^X}{(1-\hat{\pi})^{N-X} \hat{\pi}^X} \right] \xrightarrow{d} \chi_1^2,$$

with  $N = 250$ ,  $X$  the number of exceedances,  $\hat{\pi} = X/N$ . Reject the null of correct unconditional coverage if  $LR > \chi_{1,0.95}^2 = 3.841$ .

**(G.5)** Conclude: is the GARCH-VaR well-calibrated?

## Marking grid (out of 100)

| Criterion                                            | Points     |
|------------------------------------------------------|------------|
| Part A: descriptive statistics + Jarque–Bera         | 10         |
| Part B: decomposition + linear-trend fit             | 10         |
| Part C: ACF/PACF + Ljung–Box on $r_t$ and $r_t^2$    | 10         |
| Part D: ADF/KPSS + combination matrix                | 15         |
| Part E: ARMA candidates + AIC/BIC/LB                 | 15         |
| Part F: GARCH(1, 1) + standardised diagnostics + VaR | 20         |
| Part G: Kupiec back-test                             | 10         |
| Reproducibility + report quality                     | 5          |
| Oral defence                                         | 5          |
| <b>Total</b>                                         | <b>100</b> |

**Bonus.** Student- $t$  innovation distribution in the GARCH (+5 points), Christoffersen’s conditional-coverage test on top of Kupiec (+5), comparison with EGARCH or GJR-GARCH (+5).

## Submission

- **Code:** `projet_<name>.py` (or `projet_<name>.ipynb`) reproducible with seed 42.
- **Report:** 10–12 PDF pages including:
  - summary of the dataset and stylised facts;
  - all required tables and figures (one figure per part);
  - Kupiec back-test result;
  - discussion of model choice and calibration.
- **Oral defence:** 12 minutes per student, 5 minutes of questions.

### Warning

Install once: `pip install numpy scipy matplotlib pandas statsmodels arch`. The full reference pipeline runs in  $\approx 15$  seconds on a laptop CPU; the rolling Kupiec back-test (Part G) is the slowest part ( $\sim 30$  s).

*Good luck!*

Aymen Ben Brik  
aymen.benbrik@esprit.tn